

# SimpleNeuralNetworks

March 23, 2021

1 We're going to play with very simple layer ANNs. Even with just one layer we'll be able to do some fun stuff!

```
[2]: import numpy as np
      from matplotlib import pyplot
```

1.1 This is our activation function, just a simple sigmoidal curve called the logistic function


```
[3]: def logistic_func(x):
      return 1/(1+np.exp(-x))
```

Let's take a look at the logistic function by plotting in in the range of -10 to 10 using this handy built-in numpy function *linspace*. This function takes the lower bound, the upper bound, and the number of values, and returns a *linearly spaced* vector including the lower and upper bounds. **Neat!**

```
[4]: x = np.linspace(-10, 10,20)
      print(x)
```

```
[-10.          -8.94736842 -7.89473684 -6.84210526 -5.78947368
  -4.73684211 -3.68421053 -2.63157895 -1.57894737 -0.52631579
   0.52631579  1.57894737  2.63157895  3.68421053  4.73684211
   5.78947368  6.84210526  7.89473684  8.94736842 10.         ]
```

```
[5]: #Now plot it!
      pyplot.plot(x, logistic_func(x))
      pyplot.xlabel("input")
      pyplot.ylabel("output")
      pyplot.show()
```



1.2 Let's try to make a simple network that generates the following truth table

Input 1	Output
0	0
1	1

```
[6]: # Since we're only using one value, our input
# is a single value
input_vector = np.array([1])
input_weights = np.array([5])

# because we're using numpy arrays, when we do multiplication
# it will automatically perform it element-by-element
input_x_weights = input_vector * input_weights
neuron_sum = sum(input_x_weights)

# Now apply our activation function to the
# sum of the inputs*weights vector
activation = logistic_func(neuron_sum)

print(activation)
```

0.9933071490757153

```
[7]: # Since we're only using one value, our input
      # is a single value
      input_vector = np.array([0])
      input_weights = np.array([5])

      # because we're using numpy arrays, when we do multiplication
      # it will automatically perform it element-by-element
      input_x_weights = input_vector * input_weights
      neuron_sum = sum(input_x_weights)
      activation = logistic_func(neuron_sum)

      print(activation)
```

0.5

### 1.3 Hmm, we're going to have a hard time multiplying 0 by something to get a lower value...

To fix this problem, we use something called a **bias node**. Basically, we give the neural network a constant value to use for situations exactly like this. This is typically set at 1, though -1 is also a common choice.

```
[8]: # Now we're passing in a bias node with value set to 1
      # so we'll also have to give it's synapse a weight!
      input_vector = np.array([0, 1])

      # I just picked a few big values that will pull the
      # logistic function close to either 0 or 1 depending on
      # the value of the input.
      input_weights = np.array([6, -3])

      # because we're using numpy arrays, when we do multiplication
      # it will automatically perform it element-by-element
      input_x_weights = input_vector * input_weights
      neuron_sum = sum(input_x_weights)
      activation = logistic_func(neuron_sum)

      print(activation)
```

0.04742587317756678

### 1.4 We can shorten this a lot using some nice numpy built in functions!

Spend some time making sure you understand this bit of code, it's going to come back!

```
[9]: input_vector = np.array([0, 1])
      input_weights = np.array([6, -3])

      # https://docs.scipy.org/doc/numpy/reference/generated/numpy.dot.html
```

```
# numpy's dot function computes the inner product
# e.g., [a,b].[c,d] = [a*c + b*d]
activation = logistic_func(np.dot(input_vector, input_weights))
print(activation)
```

0.04742587317756678

**2 1. Try your hand at choosing weights for a network that can compute this function**

Input 1	Output
0	1
1	0

[ ]:

**3 2. Let's make it a bit more complex and add a second input!**

Input 1	Input 2	Output
0	0	0
0	1	1
1	0	1
1	1	1

[ ]:

**3.1 You should make sure your weights handle all the possible binary inputs we can give this function – like the following input vectors!**

```
[82]: input_vector = np.array([0, 0, 1])
input_vector = np.array([0, 1, 1])
input_vector = np.array([1, 0, 1])
input_vector = np.array([1, 1, 1])
```

**3.2 3. What would you have to do if we wanted to have more than one output?**

Let's try to replicate the following truth table (i.e., mirror the bits)

Input 1	Input 2	Output 1	Output 2
0	0	0	0
0	1	1	0

Input 1	Input 2	Output 1	Output 2
1	0	0	1
1	1	1	1

**3.2.1 NOTE:** Because now each input node has more than one output node to connect to, we'll need more weights.

We can store these weights as a matrix with 3 rows (one for each input plus the bias) and 2 columns (one for each output). Thus, each value is a weight from the row's input node to the column's output node.

```
[10]: input_vector = np.array([1, 0, 1])

#TODO: You'll have to change these weights!
input_weights = np.array( [[1, 1],
                           [1, 1],
                           [1, 1]])

# https://docs.scipy.org/doc/numpy/reference/generated/numpy.dot.html
# we get to take advantage of more numpy fancyness here,
# when np.dot is given a matrix, it performs matrix multiplication!
activation = logistic_func(np.dot(input_vector, input_weights))
print(np.round(activation, decimals=3))
```

[0.881 0.881]

**3.2.2** You'll want to check all the possible combinations of 2-bit inputs again, like you did in Question 3.

[ ]:

**4 4. If you're feeling brave, try your hand at this one.**

Hint: It's not as simple as it looks... You'll need more **layers** for this one!

Input 1	Input 2	Output
0	0	0
0	1	1
1	0	1
1	1	0

[ ]: